

## classify emails using LLM

```
# Import Libraries
import pandas as pd
from transformers import pipeline
from tqdm import tqdm

# Load and Prepare Dataset
df = pd.read_csv("spam.csv", encoding="latin-1")[["text", "target"]]
df = df.sample(n=100, random_state=42).reset_index(drop=True) # sample 100 for performance
df["text"] = df["text"].astype(str).str.replace(r'\s+', ' ', regex=True).str.strip()

# Load Zero-Shot Classification Pipeline
classifier = pipeline("zero-shot-classification", model="facebook/bart-large-mnli")

# Define Candidate Labels
labels = ["spam", "ham"]

# Apply Model to Classify Emails
predictions = []
for text in tqdm(df["text"]):
    result = classifier(text[:512], candidate_labels=labels)
    predictions.append(result["labels"][0]) # top label

# Store Predictions
df["predicted_label"] = predictions

# Evaluate
from sklearn.metrics import accuracy_score, classification_report

df["target"] = df["target"].str.lower()
print("Accuracy:", accuracy_score(df["target"], df["predicted_label"]))
print(classification_report(df["target"], df["predicted_label"]))
df.head()
```

config.json: 1.15k/? [00:00<00:00, 20.6kB/s]

model.safetensors: 100% 1.63G/1.63G [00:53<00:00, 46.3MB/s]

tokenizer\_config.json: 100% 26.0/26.0 [00:00<00:00, 2.08kB/s]

vocab.json: 899k/? [00:00<00:00, 5.68MB/s]

merges.txt: 456k/? [00:00<00:00, 8.50MB/s]

tokenizer.json: 1.36M/? [00:00<00:00, 17.5MB/s]

Device set to use cpu

100%|██████████| 100/100 [02:39<00:00, 1.59s/it] Accuracy: 0.73

	precision	recall	f1-score	support
ham	0.90	0.78	0.84	88
spam	0.17	0.33	0.23	12
accuracy			0.73	100
macro avg	0.54	0.56	0.53	100
weighted avg	0.81	0.73	0.76	100

	text	target	predicted_label
0	Funny fact Nobody teaches volcanoes 2 erupt, t...	ham	ham
1	I sent my scores to sophas and i had to do sec...	ham	ham
2	We know someone who you know that fancies you....	spam	spam
3	Only if you promise your getting out as SOON a...	ham	ham
4	Congratulations ur awarded either £500 of C...	spam	ham

```
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# Compute confusion matrix
cm = confusion_matrix(df["target"], df["predicted_label"], labels=["ham", "spam"])

# Display
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["ham", "spam"])
disp.plot(cmap="Blues")
```

```
plt.title("Confusion Matrix: Target vs Predicted")  
plt.show()
```

